


Core score movement | 0-100 scale



1. Executive Summary

AsterVale's reliable Month 1 Final ASV Score was 28.0523. On the matched cohort, Final ASV Score moved to 38.3724 (+10.2759 pp; 36.5738% relative).

The business gained discovery and branded recommendation, but owned citation coverage and high-intent conversion remain the binding weaknesses.

2. Reliable Month 1 Baseline

The baseline uses eligible Month 1 cells only. Category Presence was 16.9960; AI Share of Voice was 4.1962; Comparison Standing was 9.2391.

Branded Presence Quality was 67.1591, Branded Recommendation Rate was 44.3182, and Branded Prominence was 94.8864.

The baseline is reproducible from the cleaned, eligibility-filtered source data and is reported on a 0-100 scale.

3. What Genuinely Changed in Month 2

The movement comparison retains 291 cells eligible in both months and matches prompt_id, platform, and replicate; missing cells were not imputed.

Category Presence: 17.2727 -> 34.0909 (+16.8182 pp; 97.3684% relative).


Non-Branded Discovery Score: 9.0340 -> 18.0197 (+8.9857 pp; 99.4659% relative).

Branded Recommendation Rate: 43.6620 -> 59.1549 (+15.4930 pp; 35.4839% relative).


Comparison Standing: 10.6250 -> 20.0000 (+9.3750 pp; 88.2353% relative).

Relative gains are strongest where the M1 base was small. Absolute percentage-point movement is the primary interpretation; relative change is context, not a substitute for scale.

Largest prompt-family mention improvements | percentage points



Platform mention movement | percentage points



4. Where the Change Came From

Largest matched prompt-family mention improvements were hot_sleepers: mentions 6.67% -> 40.00% (+33.33 pp); recommendations 0.00% -> 20.00% (+20.00 pp)., eco_friendly: mentions 10.53% -> 42.11% (+31.58 pp); recommendations 0.00% -> 0.00% (+0.00 pp)., and care_durability: mentions 5.56% -> 27.78% (+22.22 pp); recommendations 5.56% -> 11.11% (+5.56 pp)..

The main family weaknesses were category_best: mentions 25.00% -> 37.50% (+12.50 pp); recommendations 12.50% -> 0.00% (-12.50 pp). and purchase_decision: mentions 5.26% -> 26.32% (+21.05 pp); recommendations 5.26% -> 0.00% (-5.26 pp).; both lost recommendation rate despite higher mentions.

Platform movement was led by SearchAI (+21.21 pp mentions; +8.08 pp recommendations), followed by ChatAssist (+11.00 pp mentions). No platform declined.

5. Fix Assessment

F01-F05 were classified as consistent_with_improvement because their matched target segments and at least one intended metric improved. F06 and F07 were not_evaluable_month_2 because timing or partial rollout prevents a fair Month 2 evaluation.

These are evidence-consistency judgments, not causal claims. The fix log, matched prompt-family movement, and intended metric movement are kept visible in outputs/fix_assessment.csv.

6. Remaining Binding Constraint

The clearest constraint is proof-to-recommendation conversion: Branded Recommendation Rate reached 59.1549, but Branded Owned Citation Rate remained 2.8169.

In non-branded discovery, Category Presence improved substantially, while category_best and purchase_decision recommendation rates declined. Visibility is therefore ahead of decision-ready evidence and owned proof.

7. Month 3 Priorities

1. Strengthen branded proof and owned citations: add canonical facts, supported cooling claims, care guidance, and internal links; track branded owned citation and recommendation rates.
2. Convert discovery into recommendation: improve category_best and purchase_decision content with comparison tables, value proof, and suitability qualifiers; track non-branded recommendation rate by family.
3. Improve comparison and shortlist proof: publish factual comparison criteria and trade-offs; track Comparison Standing plus brand_shortlist mention and recommendation rates.
4. Address stagnant needs: create targeted content for couples_temperature and value_under_200; track positive recommendation movement without weakening mentions.

8. Recommended Audit-Process Improvements

Keep a permanent QA ledger with raw count, deduped count, scope exclusions, non-success exclusions, incomplete five-brand runs, hierarchy failures, rank failures, and text-consistency failures by month.

Freeze the matched-cell definition before analysis and publish M1/M2 eligible-cell counts, intersection counts, and excluded-cell reasons alongside every movement claim.

Separate absolute percentage-point change from relative change, flag large relative changes from weak bases, and leave relative change undefined when M1 is zero.

Version prompt metadata, outcome coding, citation canonicalization, and fix deployment timing. Use directional GA4/GSC evidence as context only and require future audits to distinguish timing association from causal proof.

Supporting analytics context

GSC and GA4 results are directional context only. They describe search visibility and site outcomes, but do not establish that any deployed fix caused a change.

Monthly GSC non-branded impressions increased from 3,545 in 2026-06 to 4,893 in 2026-07; clicks increased from 146 to 229. Monthly GA4 sessions increased from 567 to 794 and conversions from 12 to 24.

The page-level view separates the cooling guide, comparison page, and cooling product page so Month 3 teams can monitor the content routes most relevant to the audit findings.

Fix assessment detail

F01 Entity and company-fact cleanup: consistent_with_improvement; branded presence quality and target recommendations improved in the matched evidence.

F02 Cooling Sheets Guide: consistent_with_improvement; target discovery and owned-citation signals improved in the matched evidence.

F03 Cooling-brand comparison page: consistent_with_improvement; comparison standing and target recommendation movement improved.

F04 Product claim and evidence refresh: consistent_with_improvement; branded quality and recommendation movement improved, while owned citation movement was flat.

F05 Internal-linking and hub navigation: consistent_with_improvement; category presence improved, but owned citation movement was flat.

F06 Independent expert-roundup outreach and F07 Night-sweats FAQ expansion: not_evaluable_month_2 because deployment timing or partial rollout prevented a fair Month 2 evaluation.

These classifications indicate consistency with observed movement only; they do not claim causation.

Analytical controls

The trusted dataset deduplicates run_id by latest ingested_at, keeps locked prompts and approved platforms/replicates, requires successful runs, and requires exactly one valid outcome for each of the five canonical brands.

Outcome hierarchy is validated from mentioned to shortlisted to recommended to top_choice. Mentioned brands must have unique ranks 1-5; unmentioned brands must have no rank. Structurally invalid runs are excluded.

Citation URLs are canonicalized by scheme/host/path/query rules, tracking parameters are removed, duplicate canonical citations are collapsed, and owned status is determined from the canonical brand domain.

How to read the scorecard

Absolute percentage-point movement is the primary evidence for practical change. Relative change is reported for context and can look large when the Month 1 base is weak; a zero Month 1 base has no defined relative change.

Month 1 and Month 2 movement uses only the intersection of eligible prompt_id + platform + replicate cells. Missing cells are not imputed.

The full reproducible artifacts are available in outputs/: baseline, matched comparison, prompt-family drivers, platform drivers, fix assessment, next-sprint plan, QA report, and supporting analytics summary.

Decision takeaway

AsterVale achieved meaningful Month 2 visibility gains, but the next operating challenge is converting visibility into recommendation through stronger owned proof, comparison evidence, and targeted content for stagnant high-intent segments.