


Core score movement | 0-100 scale



I. Executive Summary

AsterVale's reliable Month 1 Final ASV Score was 28.0523. On the matched cohort, Final ASV Score moved to 38.3724 (+10.2759 pp; 36.5738% relative).

The business gained discovery and branded recommendation, but owned citation coverage and high-intent conversion remain the binding weaknesses.

II. Reliable Month 1 Baseline

The baseline uses eligible Month 1 cells only. Category Presence was 16.9960; AI Share of Voice was 4.1962; Comparison Standing was 9.2391.

Branded Presence Quality was 67.1591, Branded Recommendation Rate was 44.3182, and Branded Prominence was 94.8864.

The baseline is reproducible from the cleaned, eligibility-filtered source data and is reported on a 0-100 scale.

III. What Genuinely Changed in Month 2

The movement comparison retains 291 cells eligible in both months and matches prompt_id, platform, and replicate; missing cells were not imputed.

Category Presence: 17.2727 -> 34.0909 (+16.8182 pp; 97.3684% relative).


Non-Branded Discovery Score: 9.0340 -> 18.0197 (+8.9857 pp; 99.4659% relative).

Branded Recommendation Rate: 43.6620 -> 59.1549 (+15.4930 pp; 35.4839% relative).


Comparison Standing: 10.6250 -> 20.0000 (+9.3750 pp; 88.2353% relative).

Relative gains are strongest where the M1 base was small. Absolute percentage-point movement is the primary interpretation; relative change is context, not a substitute for scale.

Largest prompt-family mention improvements | percentage points



Platform mention movement | percentage points



IV. Where the Change Came From

Largest matched prompt-family mention improvements were hot_sleepers: mentions 6.67% -> 40.00% (+33.33 pp); recommendations 0.00% -> 20.00% (+20.00 pp), eco_friendly: mentions 10.53% -> 42.11% (+31.58 pp); recommendations 0.00% -> 0.00% (+0.00 pp), and care_durability: mentions 5.56% -> 27.78% (+22.22 pp); recommendations 5.56% -> 11.11% (+5.56 pp).

The main family weaknesses were category_best: mentions 25.00% -> 37.50% (+12.50 pp); recommendations 12.50% -> 0.00% (-12.50 pp). and purchase_decision: mentions 5.26% -> 26.32% (+21.05 pp); recommendations 5.26% -> 0.00% (-5.26 pp); both lost recommendation rate despite higher mentions.

Platform movement was led by SearchAI (+21.21 pp mentions; +8.08 pp recommendations), followed by ChatAssist (+11.00 pp mentions). No platform declined.

V. Fix Assessment

F01-F05 were classified as consistent_with_improvement because their matched target segments and at least one intended metric improved. F06 and F07 were not_evaluable_month_2 because timing or partial rollout prevents a fair Month 2 evaluation.

These are evidence-consistency judgments, not causal claims. The fix log, matched prompt-family movement, and intended metric movement are kept visible in outputs/fix_assessment.csv.

VI. Remaining Binding Constraint

The clearest constraint is proof-to-recommendation conversion: Branded Recommendation Rate reached 59.1549, but Branded Owned Citation Rate remained 2.8169.

In non-branded discovery, Category Presence improved substantially, while category_best and purchase_decision recommendation rates declined. Visibility is therefore ahead of decision-ready evidence and owned proof.

VII. Month 3 Priorities

1. Strengthen branded proof and owned citations: add canonical facts, supported cooling claims, care guidance, and internal links; track branded owned citation and recommendation rates.
2. Convert discovery into recommendation: improve category_best and purchase_decision content with comparison tables, value proof, and suitability qualifiers; track non-branded recommendation rate by family.
3. Improve comparison and shortlist proof: publish factual comparison criteria and trade-offs; track Comparison Standing plus brand_shortlist mention and recommendation rates.
4. Address stagnant needs: create targeted content for couples_temperature and value_under_200; track positive recommendation movement without weakening mentions.

VIII. Recommended Audit-Process Improvements

Keep a permanent QA ledger with raw count, deduped count, scope exclusions, non-success exclusions, incomplete five-brand runs, hierarchy failures, rank failures, and text-consistency failures by month.

Freeze the matched-cell definition before analysis and publish M1/M2 eligible-cell counts, intersection counts, and excluded-cell reasons alongside every movement claim.

Separate absolute percentage-point change from relative change, flag large relative changes from weak bases, and leave relative change undefined when M1 is zero.

Version prompt metadata, outcome coding, citation canonicalization, and fix deployment timing. Use directional GA4/GSC evidence as context only and require future audits to distinguish timing association from causal proof.